

First, we need to make sure that our KinD cluster has more than one node.

Delete the existing KinD cluster (if any):

```
kind delete cluster
```

Create a YAML file like this:

```
# cluster.yaml
# kind-multinode-cluster.yaml
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
  - role: control-plane
  - role: worker
  - role: worker
```

Use the file to create a new KinD cluster with the above configuration:

```
kind create cluster --config=cluster.yaml
```

Wait for a few seconds till all the nodes are available and ready and confirm by running:

```
kubect1 get nodes
```

Create a Daemonset as follows:

```
# daemonset.yaml
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd
  labels:
    app: fluentd
spec:
  selector:
    matchLabels:
      app: fluentd
  template:
    metadata:
      labels:
        app: fluentd
    spec:
      containers:
        - name: fluentd
          image: fluent/fluentd:v1.17.1-1.0
```

Apply the manifest to the cluster:

```
kubect1 apply -f daemonset.yaml
```

Make sure that we have one pod running on each of the worker nodes:

```
kuebctl get pods -o wide
```

Limiting Daemonset to speific nodes

Label one of the nodes with gpu=true:

```
kubect1 get nodes
kubect1 label nodes kind-worker gpu=true
```

Modify the daemonset to target only nodes with this label:

```
# ... previous content
template:
  metadata:
    labels:
      app: fluentd
  spec:
    nodeSelector:
      gpu: "true"
    containers:
# ... rest of the content
```

Apply the modified daemonset:

```
kubect1 apply -f daemonset.yaml
```

Make sure that the pod is now scheduled on the designated node only:

```
kubect1 get pods -o wide
```